

## 4. Wrangling DataFrames

## Questions

- How can you select individual columns or rows from a `DataFrame` ?
- How can you subset a `DataFrame` ?
- How can you sort a `DataFrame` ?

## Objectives

- Learn how to select specific columns or rows from a `DataFrame` .
- Learn how to select rows based on conditions.
- Learn how to sort a `DataFrame` 's rows or columns.

# What is Data Wrangling?

Data wrangling is the process of cleaning, transforming, and mapping data from various sources into a format that can be used for analysis or modeling. It involves various tasks such as handling missing or duplicate data, handling outliers, and combining data from different sources.

## Subsetting, Selection and Sorting a DataFrame

- Selecting, subsetting, and sorting are essential aspects of data wrangling.
- Subsetting refers to extracting a portion of the data. For example,
  - Extract only columns 1 and 2 of the `DataFrame`
  - Give me only the rows with an even index (i.e., 2, 4, 6, ... )
- Selecting consists of extracting rows based on some condition. For example,
  - Give me all the value for which the temperature is larger than 20
  - Give me all samples that start with `ABC` , etc.
- Sorting consists of reordering data using either the index or the values of one or more columns.

# Subsetting

To illustrate subsetting, selection and sorting, we will be using the the following `DataFrame` stored in a value called `df` ,

	date mmddyy	press dbar	temp ITS-90	csal PSS-78	coxy umol/kg	ph
Sample ID						
Sample-1	40610	239.8	18.9625	35.0636	NaN	7.951
Sample-2	40610	280.7	16.1095	34.6103	192.3	NaN
Sample-3	40610	320.1	12.9729	34.2475	190.8	NaN
Sample-4	40610	341.3	11.9665	34.1884	191.3	7.780
Sample-5	40610	360.1	11.3636	34.1709	203.5	NaN
Sample-6	40610	385.0	10.4636	34.1083	193.7	NaN
Sample-7	40610	443.7	8.5897	34.0567	156.5	NaN
Sample-8	40610	497.8	7.1464	34.0424	110.7	7.496

```
In [2]: import pandas as pd
df = pd.read_csv("data/selection_dataframe.csv", index_col="Sample ID")
df.head(2)
```

```
Out[2]:
```

	date mmddyy	press dbar	temp ITS- 90	csal PSS- 78	coxy umol/kg	ph
Sample ID						
Sample-1	40610	239.8	18.9625	35.0636	NaN	7.951
Sample- 8	40610	497.8	7.1464	34.0424	110.7	7.496

# Subsetting Columns

- Selecting columns is trivial and follows the same notation as that used to with Python lists
- Selecting the `'ph'` can be written as:

```
In [3]: df['ph']
```

```
Out[3]: Sample ID  
Sample-1    7.951  
Sample-8    7.496  
Sample-2     NaN  
Sample-4    7.780  
Sample-6     NaN  
Sample-5     NaN  
Sample-7     NaN  
Sample-3     NaN  
Name: ph, dtype: float64
```



## Selecting Rows

- Subsetting on rows uses the `.loc` or `.iloc` operators
- Both `.loc` or `.iloc` require information about the rows and column.
- Example, to extract subset consisting of the row 0 and column 2, we would write:

```
df.iloc[0, 2]
```

```
In [4]: df.iloc[0, 2]
```

```
Out[4]: 18.9625
```

## Subsetting format using `iloc`

- In general, subsetting rows using `.loc` or `.iloc` can take of the following forms
- `df.iloc[row_index(es), col_index(es)]`
- As the notation indicates, we can provide a single or list of index.

Example:

```
df.iloc[1, 2]
```

```
In [5]: df.iloc[1, 2]
```

```
Out[5]: 7.1464
```

```
In [6]: df.iloc[1, [1,2,3]]
```

```
Out[6]: press dbar      497.8000  
temp ITS-90        7.1464  
csal PSS-78       34.0424  
Name: Sample-8, dtype: float64
```

```
In [7]: df.iloc[[0,4,6], [1,2,3]]
```

```
Out[7]:
```

	press dbar	temp ITS-90	csal PSS-78
Sample ID			
Sample-1	239.8	18.9625	35.0636
Sample-6	385.0	10.4636	34.1083
Sample-7	443.7	8.5897	34.0567

```
In [8]: df.loc['Sample-1', df.columns]
```

```
Out[8]: date mmddyy      40610.0000  
        press dbar      239.8000  
        temp ITS-90      18.9625  
        csal PSS-78      35.0636  
        coxy umol/kg      NaN  
        ph              7.9510  
        Name: Sample-1, dtype: float64
```

## Subsetting format using `loc`

- in general subsetting rows using `.loc` is written as
- `df.loc[row_label(s), column_label(es)]`
- As the notation indicates, we can provide a single or list of labels



INPUT

```
df.loc['Sample-1', df.columns]
```

Output

```
date mddyy      40610.0000
press dbar      239.8000
temp ITS-90     18.9625
csal PSS-78     35.0636
coxy umol/kg    NaN
ph              7.9510
Name: Sample-1, dtype: float64
```

## The `:` Operator In Python (Refresher)

- In python the `:` operator is used for slicing ordered collections (e.g. lists).
- "`:`" can optionally take left and right operators:
  - `X:Y` means from X to and not including Y
  - `X:` means from X to the end of the list
  - `:Y` means from the beginning all the way to and not including Y
  - `:` the entire collection

```
In [9]: x = ['a', 'b', 'c', 'd', 'e']  
        x[1:4]
```

```
Out[9]: ['b', 'c', 'd']
```

```
In [10]: x = ['a', 'b', 'c', 'd', 'e']  
         x[1:]
```

```
Out[10]: ['b', 'c', 'd', 'e']
```

```
In [11]: x = ['a', 'b', 'c', 'd', 'e']  
x[:]
```

```
Out[11]: ['a', 'b', 'c', 'd', 'e']
```

```
In [12]: df.columns
```

```
Out[12]: Index(['date mddyy', 'press dbar', 'temp ITS-90', 'csal PSS-7  
8',  
               'coxy umol/kg', 'ph'],  
              dtype='object')
```

## Indexing Using Slicing operator

INPUT

```
df.loc['Sample-1', :]
```

Output

```
date mddyy      40610.0000
press dbar      239.8000
temp ITS-90      18.9625
csal PSS-78      35.0636
coxy umol/kg      NaN
ph              7.9510
Name: Sample-1, dtype: float64
```

```
In [13]: df.loc['Sample-1', :]
```

```
Out[13]: date mmddyy      40610.0000  
         press dbar      239.8000  
         temp ITS-90      18.9625  
         csal PSS-78      35.0636  
         coxy umol/kg      NaN  
         ph              7.9510  
         Name: Sample-1, dtype: float64
```



## Difference Between `.loc` and `.iloc`

Both used to subset of a `DataFrame`.

- `.loc` relies on the names of the indexes and headers
- `.iloc` relies instead on the index and header number
- `df.loc['Sample-1', ["date mmddyy", "ph"]]` means row index by the label 'Sample-1' and columns "date mmddyy", "ph"
- `df.iloc[0, [0,1,2]]` means row 0 and columns 0, 1 and 2
- You cannot mix indexes and labels and the following is incorrect
  - `df.iloc[0, ["date mmddyy", "ph"]]` returns error
  - `df.iloc[0, df.columns]` returns error

```
In [14]: df.loc['Sample-1', ["date mmddyy", "ph"]]
```

```
Out[14]: date mmddyy    40610.000  
         ph            7.951  
         Name: Sample-1, dtype: float64
```

```
In [15]: df.iloc[0, [0,1,2]]
```

```
Out[15]: date mddyy      40610.0000  
         press dbar      239.8000  
         temp ITS-90      18.9625  
         Name: Sample-1, dtype: float64
```

```
In [16]: ### The following should return an error.  
df.iloc[0, df.columns]
```

```
-----  
-----  
IndexError                                Traceback (most recent  
call last)
```

```
Cell In[16], line 2
```

```
      1 ### The following should return an error.  
----> 2 df.iloc[0, df.columns]
```

```
File ~/miniconda3/envs/temp/lib/python3.9/site-packages/pandas/  
core/indexing.py:1067, in _LocationIndexer._getitem__(self, ke  
y)
```

```
    1065     if self._is_scalar_access(key):  
    1066         return self.obj._get_value(*key, takeable=self.  
_takeable)  
-> 1067     return self._getitem_tuple(key)  
    1068 else:  
    1069     # we by definition only have the 0th axis  
    1070     axis = self.axis or 0
```

```
File ~/miniconda3/envs/temp/lib/python3.9/site-packages/pandas/  
core/indexing.py:1563, in _iLocIndexer._getitem_tuple(self, tu  
ple)
```

```
    1561 def _getitem_tuple(self, tup: tuple):  
-> 1563     tup = self._validate_tuple_indexer(tup)  
    1564     with suppress(IndexingError):  
    1565         return self._getitem_lowerdim(tup)
```

```
File ~/miniconda3/envs/temp/lib/python3.9/site-packages/pandas/  
core/indexing.py:873, in _LocationIndexer._validate_tuple_index  
er(self, key)
```

```
    871 for i, k in enumerate(key):  
    872     try:  
--> 873         self._validate_key(k, i)  
    874     except ValueError as err:  
    875         raise ValueError(  
    876     "iLocation by tuple only supports integer location"
```

```
877         f"[{self._valid_types}] types"
878     ) from err
```

File ~/miniconda3/envs/temp/lib/python3.9/site-packages/pandas/core/indexing.py:1477, in `_iLocIndexer._validate_key(self, key, axis)`

```
1475 # check that the key has a numeric dtype
1476 if not is_numeric_dtype(arr.dtype):
-> 1477     raise IndexError(f".iloc requires numeric indexers,
got {arr}")
1479 # check that the key does not exceed the maximum size o
f the index
1480 if len(arr) and (arr.max() >= len_axis or arr.min() < -
len_axis):
```

`IndexError: .iloc requires numeric indexers, got ['date mmddyy' 'press dbar' 'temp ITS-90' 'csal PSS-78' 'coxy umol/kg' 'ph']`

## Selecting using Conditional Expressions

In Pandas, comparison operations (" $<$ ", " $>$ ", " $==$ ", " $>=$ ", " $<=$ ", " $!=$ ") works in a similar as in Python

- An applied series, then a series of booleans (True or False) are returned

	date mmdyy	press dbar	temp ITS-90	csal PSS-78	coxy umol/kg	ph
Sample ID						
Sample-1	40610	239.8	18.9625	35.0636	NaN	7.951
Sample-2	40610	280.7	16.1095	34.6103	192.3	NaN
Sample-3	40610	320.1	12.9729	34.2475	190.8	NaN
Sample-4	40610	341.3	11.9665	34.1884	191.3	7.780
Sample-5	40610	360.1	11.3636	34.1709	203.5	NaN
Sample-6	40610	385.0	10.4636	34.1083	193.7	NaN
Sample-7	40610	443.7	8.5897	34.0567	156.5	NaN
Sample-8	40610	497.8	7.1464	34.0424	110.7	7.496

- What does the following return?

```
df['press dbar'] < 380
```

```
In [17]: df['press dbar'] < 380
```

```
Out[17]: Sample ID
Sample-1    True
Sample-8    False
Sample-2     True
Sample-4     True
Sample-6    False
Sample-5     True
Sample-7    False
Sample-3     True
Name: press dbar, dtype: bool
```

## Subsetting Based on Conditional Expressions

- It turns out that we can also subset a `DataFrame` using a list of booleans. For example:

INPUT

```
df[[True, True, False, False, False, False, False, False]]
```

OUTPUT

date mmddyy	press dbar	temp ITS-90	csal PSS-78	coxy
umol/kg ph				
Sample ID				
Sample-1	40610	239.8	18.9625 35.0636	NaN 7.951
Sample-2	40610	280.7	16.1095 34.6103	192.3 NaN




In [18]: `df[[True, True, False, False, False, False, False, False]]`

Out[18]:

	date mmddyy	press dbar	temp ITS- 90	csal PSS- 78	coxy umol/kg	ph
Sample ID						
Sample-1	40610	239.8	18.9625	35.0636	NaN	7.951
Sample-8	40610	497.8	7.1464	34.0424	110.7	7.496

- The python input means return the first two row (associated with True) and ignore the lines associated with False



- Note that in when subsetting using Booleans, we don't need to use either `.loc` or `.iloc`

## Subsetting Based on Conditional Expressions - Cont'd

- Conditional expression can be used directly in a `DataFrame`

INPUT

```
df[ df['press dbar'] < 380 ]
```

OUTPUT

date mmdyy	press dbar	temp ITS-90	csal PSS-78	coxy
umol/kg ph				
Sample ID				
Sample-1	40610 239.8	18.9625 35.0636	NaN	7.951
Sample-2	40610 280.7	16.1095 34.6103	192.3	NaN
Sample-3	40610 320.1	12.9729 34.2475	190.8	NaN
Sample-4	40610 341.3	11.9665 34.1884	191.3	7.780
Sample-5	40610 360.1	11.3636 34.1709	203.5	NaN

```
In [19]: df[ df['press dbar'] < 380 ]
```

Out[19]:

	date mmddyy	press dbar	temp ITS- 90	csal PSS- 78	coxy umol/kg	ph
Sample ID						
Sample-1	40610	239.8	18.9625	35.0636	NaN	7.951
Sample-2	40610	280.7	16.1095	34.6103	192.3	NaN
Sample-4	40610	341.3	11.9665	34.1884	191.3	7.780
Sample-5	40610	360.1	11.3636	34.1709	203.5	NaN
Sample-3	40610	320.1	12.9729	34.2475	190.8	NaN

# Sorting

- Sorting is important aspect of data wrangling. For example:
  - In data analysis: sorting is used to identify patterns, trends, and outliers in the data.
  - in data cleaning, sorting can help to identify and remove duplicate entries, missing values, and other errors in the data.
  - etc.
- Data can sorted on the index using `sort_index()`
- Data can sorted on of the columns using `sort_values()`
  - Can take a single column or a list of columns to sort
- By default index and vlaues are sorted in ascending order but the behavior can be changed by setting `ascending=True` as parameter to the function.

In [20]: `df.sort_index()`

	date mmddyy	press dbar	temp ITS- 90	csal PSS- 78	coxy umol/kg	ph
Sample ID						
Sample-1	40610	239.8	18.9625	35.0636	NaN	7.951
Sample-2	40610	280.7	16.1095	34.6103	192.3	NaN
Sample-3	40610	320.1	12.9729	34.2475	190.8	NaN
Sample-4	40610	341.3	11.9665	34.1884	191.3	7.780
Sample-5	40610	360.1	11.3636	34.1709	203.5	NaN
Sample-6	40610	385.0	10.4636	34.1083	193.7	NaN
Sample-7	40610	443.7	8.5897	34.0567	156.5	NaN
Sample-8	40610	497.8	7.1464	34.0424	110.7	7.496

```
In [21]: df.sort_values(by='press dbar')
```

	date mmddyy	press dbar	temp ITS- 90	csal PSS- 78	coxy umol/kg	ph
Sample ID						
Sample-1	40610	239.8	18.9625	35.0636	NaN	7.951
Sample-2	40610	280.7	16.1095	34.6103	192.3	NaN
Sample-3	40610	320.1	12.9729	34.2475	190.8	NaN
Sample-4	40610	341.3	11.9665	34.1884	191.3	7.780
Sample-5	40610	360.1	11.3636	34.1709	203.5	NaN
Sample-6	40610	385.0	10.4636	34.1083	193.7	NaN
Sample-7	40610	443.7	8.5897	34.0567	156.5	NaN
Sample-8	40610	497.8	7.1464	34.0424	110.7	7.496

```
In [22]: df.sort_values(by='press dbar', ascending=False)
```

	date mmddyy	press dbar	temp ITS- 90	csal PSS- 78	coxy umol/kg	ph
Sample ID						
Sample-8	40610	497.8	7.1464	34.0424	110.7	7.496
Sample-7	40610	443.7	8.5897	34.0567	156.5	NaN
Sample-6	40610	385.0	10.4636	34.1083	193.7	NaN
Sample-5	40610	360.1	11.3636	34.1709	203.5	NaN
Sample-4	40610	341.3	11.9665	34.1884	191.3	7.780
Sample-3	40610	320.1	12.9729	34.2475	190.8	NaN
Sample-2	40610	280.7	16.1095	34.6103	192.3	NaN
Sample-1	40610	239.8	18.9625	35.0636	NaN	7.951



```
In [23]: df.sort_values(by=['ph', 'press dbar'], ascending=False)
```

	date mmddyy	press dbar	temp ITS- 90	csal PSS- 78	coxy umol/kg	ph
Sample ID						
Sample-1	40610	239.8	18.9625	35.0636	NaN	7.951
Sample-4	40610	341.3	11.9665	34.1884	191.3	7.780
Sample-8	40610	497.8	7.1464	34.0424	110.7	7.496
Sample-7	40610	443.7	8.5897	34.0567	156.5	NaN
Sample-6	40610	385.0	10.4636	34.1083	193.7	NaN
Sample-5	40610	360.1	11.3636	34.1709	203.5	NaN
Sample-3	40610	320.1	12.9729	34.2475	190.8	NaN
Sample-2	40610	280.7	16.1095	34.6103	192.3	NaN

## Key Points

- Select columns by using `[ "column name" ]` or rows by using the `loc` attribute.
- Sort based on values in a column by using the `sort_values` method.

# Exercise

Try it yourself! Going back to our `20_sales_records.xlsx` file, identify which orders are `Online` and `High Priority`

- Read the first couple rows to get a sense of the data. Which column reflects `Online` or `Offline` Status.
- A `High Priority` order is denoted by `H` in one of the columns. Identify which column.
- HINT: Use the `loc` method

